

NLP_aggressive_classifier_final

July 10, 2020

1 Implementation

- Purpose of this project is to analyse the Jigsaw comment dataset and classify the data based on toxicity
- Three algorithms were tested based on the Naive Bayes, Logistic regression and support vector machine for different values of n grams and TFIDF norm values.
- Appropriate value of 'n' in n-grams is selected using the result from above algorithms also better norm for TFIDF is selected.
- Algorithms giving good performance are taken as baseline
- Finally hyperparameters of the algorithms are tuned for better results.

2 Data

Jigsaw toxic comment data is chosen from kaggle competition for training and testing purpose. Jigsaw toxic comment data [2] is a set of data containing toxic comments and binary indicators to represent whether the data is toxic or non-toxic also level of toxicity in the message is also presented.

Comments are in five categories: - severe_toxic

- obscene

- threat - insult

- identity_hate

3 Loading necessary dependency for operation

```
[1]: import nltk
import pickle
import re
import time
import warnings
import csv
import requests

import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```

import itertools
import traceback
from wordcloud import WordCloud

from sklearn import preprocessing
from nltk.stem.porter import *

from sklearn import metrics
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfTransformer
# from performance import *
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import SGDClassifier
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import KFold
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, confusion_matrix
from collections import Counter

```

```

[2]: def telegram_bot_sendtext(bot_message):

    bot_token = ''
    bot_chatID = ''
    send_text = 'https://api.telegram.org/bot' + bot_token + '/sendMessage?
    ↪chat_id=' + bot_chatID + '&parse_mode=Markdown&text=' + bot_message

    response = requests.get(send_text)

    bot_token_alan = ''
    bot_chatID_alan = ''
    send_text_alan = 'https://api.telegram.org/bot' + bot_token_alan + '/
    ↪sendMessage?chat_id=' + bot_chatID_alan + '&parse_mode=Markdown&text=' +
    ↪bot_message

    response = requests.get(send_text_alan)

    return response.json()

```

```

[3]: print(pickle.compatible_formats)

```

```
['1.0', '1.1', '1.2', '1.3', '2.0', '3.0', '4.0']
```

4 Exploratory data analysis

4.1 Read the dataset

```
[5]: train_data = pd.read_csv("../data/train.csv", index_col=False)
test_data = pd.read_csv("../data/test.csv", index_col=False)
test_data_labels = pd.read_csv("../data/test_labels.csv", index_col=False)
test_data_labels.head()
print(type(test_data))
```

```
<class 'pandas.core.frame.DataFrame'>
```

4.2 Preprocess the testing dataset

- If we look at the test dataset it has got three categories -1,0 and 1.
- As per the guidelines from the kaggle it is stated that column value with -1 is not taken for grading.
- So we need to remove that column since we actually don't know to which categories(toxic or non toxic) the data belongs to.
- By removing the appropriate rows we can compare the performance for prediction and ground truth value.
- Since -1 value is not consistent for different columns of toxicity we need to make local copy and remove appropriate rows.

```
[6]: test_data_toxic = test_data[['comment_text']].copy()
test_data_labels_toxic = test_data_labels['toxic'].copy().to_frame('toxic')
x = list(np.where(test_data_labels_toxic['toxic'] == -1)[0])
test_data_labels_toxic = test_data_labels_toxic.drop(test_data_labels_toxic.
→index[x])
test_data_toxic = test_data_toxic.drop(test_data_toxic.index[x])
print(len(test_data_toxic), len(test_data_labels_toxic))
```

```
63978 63978
```

4.3 Dataset size

```
[7]: print("Train data size:", train_data.shape)
print("Test data size: ", test_data.shape)
```

```
Train data size: (159571, 8)
```

```
Test data size: (153164, 2)
```

4.4 Charactersitics of data

4.4.1 Train data

```
[ ]: train_data.describe()
```

```
[ ]:
      toxic      severe_toxic  ...      insult  identity_hate
count  159571.000000  159571.000000  ...  159571.000000  159571.000000
mean      0.095844      0.009996  ...      0.049364      0.008805
std      0.294379      0.099477  ...      0.216627      0.093420
min      0.000000      0.000000  ...      0.000000      0.000000
25%      0.000000      0.000000  ...      0.000000      0.000000
50%      0.000000      0.000000  ...      0.000000      0.000000
75%      0.000000      0.000000  ...      0.000000      0.000000
max      1.000000      1.000000  ...      1.000000      1.000000
```

[8 rows x 6 columns]

```
[ ]: train_data.groupby('toxic').describe()
```

```
[ ]:
      severe_toxic
      count      mean      std  min  25%  ...      identity_hate
75%  max
toxic
0      144277.0  0.000000  0.000000  0.0  0.0  ...      0.0  0.0  0.0
0.0  1.0
1      15294.0  0.104289  0.305645  0.0  0.0  ...      0.0  0.0  0.0
0.0  1.0
```

[2 rows x 40 columns]

4.4.2 Test data

```
[ ]: test_data.describe()
```

```
[ ]:
      id      comment_text
count      153164      153164
unique      153164      153164
top      e8696ec7f26006e5 == Ham's wife == \n\n Why not bring it up on t...
freq      1      1
```

```
[ ]:
```

4.5 Visualize toxic comments counts

```
[ ]: sns.set(style="darkgrid")

f, axes = plt.subplots(1, 3, figsize=(15, 5), sharex=True)

graph_1 = sns.countplot(x = 'toxic', data = train_data, palette = 'hls',
    ↳ax=axes[0], order = [1, 0])
graph_1.set_title('toxic count')

graph_2 = sns.countplot(x = 'severe_toxic', data = train_data, palette = 'hls',
    ↳ax=axes[1], order = [1, 0])
graph_2.set_title('severe_toxic count')

graph_3 = sns.countplot(x = 'obscene', data = train_data, palette = 'hls',
    ↳ax=axes[2], order = [1, 0])
graph_3.set_title('obscene count')

f.subplots_adjust(wspace=0.4)

f, axes = plt.subplots(1, 3, figsize=(15, 5), sharex=True)

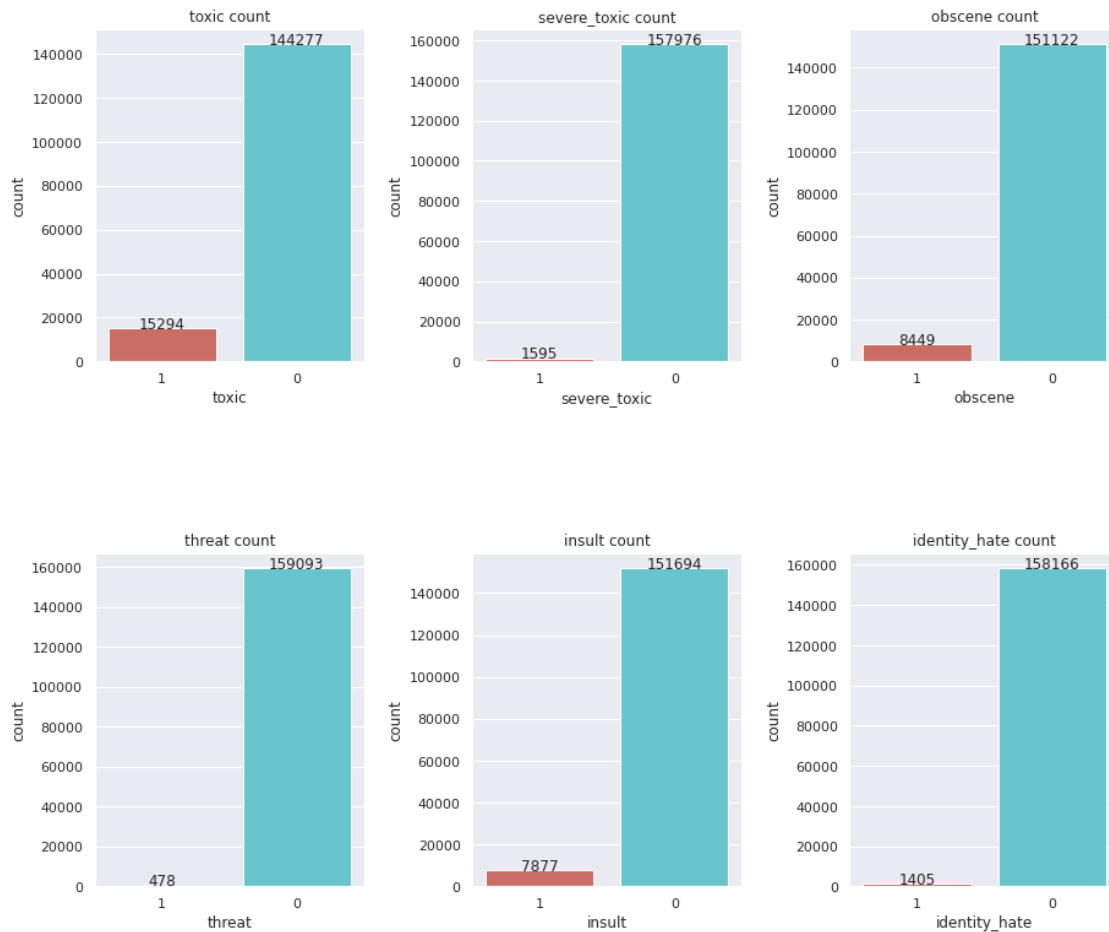
graph_4 = sns.countplot(x = 'threat', data = train_data, palette = 'hls',
    ↳ax=axes[0], order = [1, 0])
graph_4.set_title('threat count')

graph_5 = sns.countplot(x = 'insult', data = train_data, palette = 'hls',
    ↳ax=axes[1], order = [1, 0])
graph_5.set_title('insult count')

graph_6 = sns.countplot(x = 'identity_hate', data = train_data, palette =
    ↳'hls', ax=axes[2], order = [1, 0])
graph_6.set_title('identity_hate count')

f.subplots_adjust(wspace=0.4)

for graph in [graph_1, graph_2, graph_3, graph_4, graph_5, graph_6]:
    for p in graph.patches:
        height = p.get_height()
        graph.text(p.get_x()+p.get_width()/2., height + 0.1,height ,ha="center")
```



```
[8]: text = train_data['comment_text'].tolist()
tags = train_data['toxic'].tolist()

# Generate strings for wordcloud
toxic_text = " ".join([words for words, tag in zip(text, tags) if tag == 1])
notoxic_text = " ".join([words for words, tag in zip(text, tags) if tag == 0])

toxic_wordcloud = WordCloud(height=500, width=500, background_color='white',
    ↳ collocations=False).generate(toxic_text)
notoxic_wordcloud = WordCloud(height=500, width=500, background_color='white',
    ↳ collocations=False).generate(notoxic_text)

# Save wordclouds
toxic_wordcloud.to_file("../images/Final/wordcloud/unigram_toxic_wordcloud.png")
notoxic_wordcloud.to_file("../images/Final/wordcloud/unigram_notoxic_wordcloud.
    ↳ png")
```

```
[8]: <wordcloud.wordcloud.WordCloud at 0x7ff7f754e780>
```

4.5.1 Words used in toxic sentences

4.5.2 Words used in no toxic sentences

4.6 Comments length based on number of words used

```
[ ]: train_data["length"] = train_data["comment_text"].apply(len)
train_data.head(n=10)
```

```
[ ]:
      id ... length
0  0000997932d777bf ...    264
1  000103f0d9cfb60f ...    112
2  000113f07ec002fd ...    233
3  0001b41b1c6bb37e ...    622
4  0001d958c54c6e35 ...     67
5  00025465d4725e87 ...     65
6  0002bcb3da6cb337 ...     44
7  00031b1e95af7921 ...    115
8  00037261f536c51d ...    472
9  00040093b2687caa ...     70
```

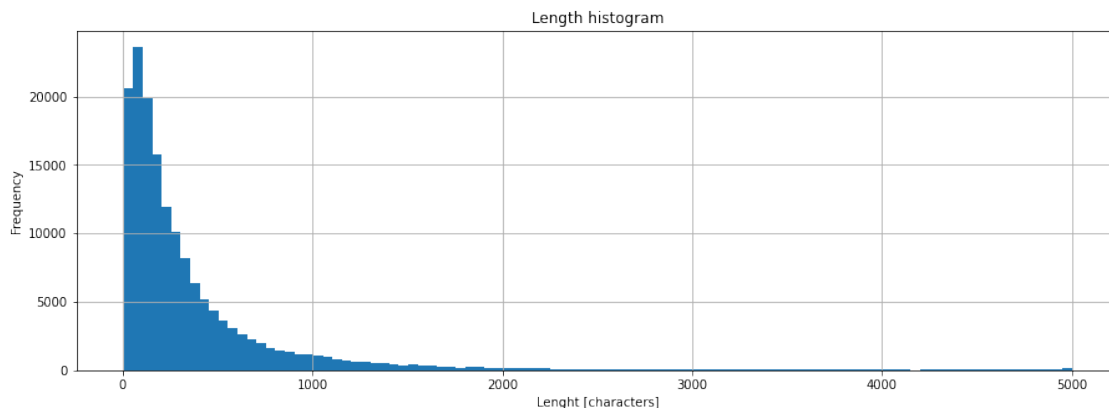
[10 rows x 9 columns]

4.6.1 Visualize the counts using histogram

```
[ ]: f = plt.figure(figsize=(15, 5))

plt.hist(train_data["length"], 100)
plt.title('Length histogram')
plt.xlabel('Length [characters]')
plt.ylabel('Frequency')

plt.grid()
plt.show()
```



4.6.2 Charactersitics of length of data

```
[ ]: train_data.length.describe()
```

```
[ ]: count    159571.000000
      mean      394.073221
      std       590.720282
      min        6.000000
      25%       96.000000
      50%      205.000000
      75%      435.000000
      max     5000.000000
      Name: length, dtype: float64
```

5 Performance comparison of algorithms w.r.t diferrent features

5.1 Extract the class labels

```
[ ]: X_train = train_data['comment_text'].values
      # X_train = X_train[0:300]
      Y_train = train_data['toxic'].values
      # Y_train = Y_train[0:300]
      print(type(X_train))
      X_test = test_data_toxic['comment_text'].values
      Y_test = test_data_labels_toxic['toxic'].values
```

```
<class 'numpy.ndarray'>
```

5.2 Train Naive Bayes algorithm iteratively for different sets of features

```
[ ]: def naive_bayes():
      nb_clf = Pipeline([
          ('vect', CountVectorizer()),
          ('tfidf', TfidfTransformer()),
          ('clf', MultinomialNB())
      ])

      param_grid = {'vect__ngram_range': ((1, 1), (1, 2), (1, 3)),
                    'tfidf__norm': ('l1', 'l2')}

      kfold = KFold(n_splits=10, random_state=23)
      gs_nb_clf = GridSearchCV(nb_clf, param_grid=param_grid, n_jobs=-1, cv=kfold,
      ↪scoring='accuracy')
```



```

start = time.time()
gs_nb_clf = gs_nb_clf.fit(X_train, Y_train)

end = time.time()
print("Time taken:", end - start)
naive_bayes = "Execution of gridsearch for naive bayes is complete. Execution_
→time: "+str(end - start)+" sec"
with open("../model/Grid_search_modes/gs_nb_features.pkl", 'wb') as_
→gs_nb_clf_file:
    pickle.dump(gs_nb_clf, gs_nb_clf_file, pickle.HIGHEST_PROTOCOL)
return naive_bayes

var = None
naive_bayes = None
try:
    naive_bayes = naive_bayes()
except:
    var = traceback.format_exc()

if var == None:
    telegram_bot_sendtext(naive_bayes)
else:
    telegram_bot_sendtext(var)

```

/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_split.py:296:
FutureWarning: Setting a random_state has no effect since shuffle is False. This
will raise an error in 0.24. You should leave random_state to its default
(None), or set shuffle=True.

FutureWarning
/usr/local/lib/python3.6/dist-
packages/joblib/externals/loky/process_executor.py:691: UserWarning: A worker
stopped while some jobs were given to the executor. This can be caused by a too
short worker timeout or by a memory leak.

"timeout or by a memory leak.", UserWarning

Time taken: 2531.117730617523

5.3 Train Logistic regression algorithm iteratively for different sets of features

```

[ ]: lr_clf = Pipeline([
    ('vect', CountVectorizer()),
    ('tfidf', TfidfTransformer()),
    ('classifier', LogisticRegression()),
])

param_grid = {'vect__ngram_range': ((1, 1), (1, 2), (1, 3)),
              'tfidf__norm': ('l1', 'l2')}

```

```

    }

kfold = KFold(n_splits=10, random_state=23)
gs_lr_clf = GridSearchCV(lr_clf, param_grid=param_grid, n_jobs=-1, cv=kfold,
    ↳scoring='accuracy')

start = time.time()
gs_lr_clf = gs_lr_clf.fit(X_train, Y_train)
end = time.time()
print("Time taken:", end - start)

logistic_regression = "Execution of gridsearch for logistic_regression is
    ↳complete. Execution time: "+str(end - start)+" sec"
telegram_bot_sendtext(logistic_regression)

with open("../model/Grid_search_modes/gs_lr_features.pkl", 'wb') as
    ↳gs_lr_clf_file:
    pickle.dump(gs_lr_clf, gs_lr_clf_file, pickle.HIGHEST_PROTOCOL)

```

/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_split.py:296:
FutureWarning: Setting a random_state has no effect since shuffle is False. This
will raise an error in 0.24. You should leave random_state to its default
(None), or set shuffle=True.

FutureWarning
/usr/local/lib/python3.6/dist-
packages/joblib/externals/loky/process_executor.py:691: UserWarning: A worker
stopped while some jobs were given to the executor. This can be caused by a too
short worker timeout or by a memory leak.

"timeout or by a memory leak.", UserWarning
/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/_logistic.py:940:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

Time taken: 4871.411665439606

5.4 Train Support Vector Machine iteratively for different sets of features

```

[ ]: svm_clf = Pipeline([
    ('vect', CountVectorizer()),
    ('tfidf', TfidfTransformer()),

```

```

        ('clf', SGDClassifier(tol=None, max_iter=100))
    ])

    param_grid = {'vect_ngram_range': ((1, 1), (1, 2), (1, 3)),
                  'tfidf__norm': ('l1', 'l2')}

    kfold = KFold(n_splits=10, random_state=23)
    gs_svm_clf = GridSearchCV(svm_clf, param_grid=param_grid, n_jobs=-1, cv=kfold,
                              scoring='accuracy')

    start = time.time()
    gs_svm_clf = gs_svm_clf.fit(X_train, Y_train)
    # gs_svm_clf = svm_clf.fit(X_train, Y_train)
    end = time.time()
    print("Time taken:", end - start)

    svm = "Execution of gridsearch for support vector machine is complete.
    Execution time: "+str(end - start)+" sec"
    telegram_bot_sendtext(svm)

    with open("../model/Grid_search_modes/gs_svm_features.pkl", 'wb') as f:
        gs_svm_clf_file:
        pickle.dump(gs_svm_clf, gs_svm_clf_file, pickle.HIGHEST_PROTOCOL)

```

/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_split.py:296:
FutureWarning: Setting a random_state has no effect since shuffle is False. This
will raise an error in 0.24. You should leave random_state to its default
(None), or set shuffle=True.

FutureWarning

Time taken: 3118.977990627289

5.5 Plot the graph

```

[11]: def draw_performance_comparison(x, y):
        fig = plt.figure(figsize=(12,6))
        ax1 = fig.add_subplot(111)

        ax1.plot(x, y[0], label="Naive Bayes")
        ax1.plot(x, y[1], label="Logistic Regression")
        ax1.plot(x, y[2], label="Support Vector Machine")

        plt.xlabel('Features')
        plt.ylabel('Validation Accuracy')
        plt.title('Performance Comparison of Algorithms w.r.t different Features')
        ax1.legend(loc=2)


```

```
plt.grid(True)
plt.savefig("../images/Final/Performance_graph/performance.png")
plt.savefig("../images/Final/Performance_graph/performance.pdf")
plt.show()
```

```
[12]: with open("../model/Grid_search_modes/gs_nb_features.pkl", 'rb') as f:
    gs_nb_clf_file = f
    gs_nb_clf = pickle.load(gs_nb_clf_file)
with open("../model/Grid_search_modes/gs_lr_features.pkl", 'rb') as f:
    gs_lr_clf_file = f
    gs_lr_clf = pickle.load(gs_lr_clf_file)
with open("../model/Grid_search_modes/gs_svm_features.pkl", 'rb') as f:
    gs_svm_clf_file = f
    gs_svm_clf = pickle.load(gs_svm_clf_file)

x = list()
params = gs_nb_clf.cv_results_['params']
for param in params:
    norm = "TDIDF norm: " + str(param['tfidf__norm'])
    ngram_range = "N-gram range" + str(param['vect__ngram_range'])
    x.append(norm + "\n" + ngram_range)

nb_means = gs_nb_clf.cv_results_['mean_test_score']
lr_means = gs_lr_clf.cv_results_['mean_test_score']
svm_means = gs_svm_clf.cv_results_['mean_test_score']
y = [nb_means, lr_means, svm_means]
performance = [{"Naive_bayes:", y[0]}, {"Logistic regression:", y[1]}, {"Support_
    vector machine:", y[2]}]
telegram_bot_sendtext(str(performance))
draw_performance_comparison(x, y)
```



6 Hyperparameter Tuning

6.0.1 Tune hyperparameters of Naive Bayes

```
[ ]: estimators_nb = [
    ('vect', CountVectorizer(ngram_range=(1,1))),
    ('tfidf', TfidfTransformer(norm='l2')),
    ('clf', MultinomialNB())
]
nb_clf = Pipeline(estimators_nb)

param_grid = dict(
    clf__alpha = [0.001, 0.01, 0.1, 1, 10],
)
kfold = KFold(n_splits=10, random_state=23)
gs_nb_clf = GridSearchCV(nb_clf, param_grid=param_grid, n_jobs=-1, cv=kfold,
    ↳scoring='accuracy')

start = time.time()
gs_nb_clf = gs_nb_clf.fit(X_train, Y_train)
end = time.time()
print("Time taken:", end - start)

naive_bayes_tuned = "Execution of hyperparameter tuning for naive bayes is
    ↳complete. Execution time: "+str(end - start)+" sec"
telegram_bot_sendtext(naive_bayes_tuned)
```

```
with open("../model/Grid_search_modes/gs_nb_tuned.pkl", 'wb') as gs_nb_clf_file:
    pickle.dump(gs_nb_clf, gs_nb_clf_file, pickle.HIGHEST_PROTOCOL)
```

/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_split.py:296:
FutureWarning: Setting a random_state has no effect since shuffle is False. This
will raise an error in 0.24. You should leave random_state to its default
(None), or set shuffle=True.

FutureWarning

Time taken: 415.47970819473267

6.1 Performance

```
[ ]: def draw_hp_performance_nb(x, y):
    fig = plt.figure(figsize=(12,6))
    ax1 = fig.add_subplot(111)

    ax1.plot(x, y, label="Naive Bayes")
    ax1.annotate(str(round(max(y), 4)), xy=(x[y.index(max(y))], y[y.
→index(max(y))]),
                xytext=(x[y.index(max(y))], max(y) - 0.01),
                arrowprops=dict(facecolor='black', shrink=0.05))
    plt.xlabel('Hyperparameters')
    plt.ylabel('Validation Accuracy')
    plt.title('Result of Naive Bayes for different hyperparameter values')
    plt.grid(True)
    plt.savefig("../images/Final/naive_bayes_hyperparameter_tuning/nb.
→png")
    plt.savefig("../images/Final/naive_bayes_hyperparameter_tuning/nb.pdf")
    plt.show()
```

```
[ ]: # NB Results
with open("../model/Grid_search_modes/gs_nb_tuned.pkl", 'rb') as gs_nb_clf_file:
    gs_nb_clf = pickle.load(gs_nb_clf_file)

with open("../model/Grid_search_modes/gs_nb_features.pkl", 'rb') as _
→gs_nb_features_file:
    gs_nb_features = pickle.load(gs_nb_features_file)

x = list()
params = gs_nb_clf.cv_results_['params']
for param in params:
    alpha = "alpha: " + str(param['clf__alpha'])
    x.append(alpha)


nb_means = gs_nb_clf.cv_results_['mean_test_score']
y = nb_means
```

```

draw_hp_performance_nb(x, y)
naive_bayes_tuned = str("Best validation accuracy for naive bayes: %f using %s" %
    ↳ % (max(y), x[y.index(max(y))]))

telegram_bot_sendtext(naive_bayes_tuned)
print("Best: %f using %s" % (max(y), x[y.index(max(y))]))

```



Best: 0.946939 using alpha: 0.01

6.2 Tune hyperparameters of Logistic regressor

```

[ ]: estimators_lr = [
    ('vect', CountVectorizer(ngram_range=(1,1))),
    ('tfidf', TfidfTransformer(norm='l2')),
    ('clf', LogisticRegression(class_weight='balanced'))
]
lr_clf = Pipeline(estimators_lr)

param_grid = dict(
    clf__C = [0.01, 0.1, 1, 10, 100],
    clf__solver = ['newton-cg', 'liblinear', 'lbfgs']
)
kfold = KFold(n_splits=10, random_state=23)
gs_lr_clf = GridSearchCV(lr_clf, param_grid=param_grid, n_jobs=-1, cv=kfold,
    ↳ scoring='accuracy')

```

```

start = time.time()
with warnings.catch_warnings():
    warnings.filterwarnings("ignore")
    gs_lr_clf = gs_lr_clf.fit(X_train, Y_train)
end = time.time()
print("Time taken:", end - start)

lr_tuned = "Execution of hyperparameter tuning for logistic regression is_
↪complete. Execution time: "+str(end - start)+" sec"
telegram_bot_sendtext(lr_tuned)

with open("../model/Grid_search_modes/gs_lr_tuned.pkl", 'wb') as gs_lr_clf_file:
    pickle.dump(gs_lr_clf, gs_lr_clf_file, pickle.HIGHEST_PROTOCOL)

```

/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_split.py:296:
FutureWarning: Setting a random_state has no effect since shuffle is False. This
will raise an error in 0.24. You should leave random_state to its default
(None), or set shuffle=True.

FutureWarning

Time taken: 2114.2145347595215

6.3 Performance

```

[ ]: def draw_hp_performance_lr(x, y):
    fig = plt.figure(figsize=(20,9))
    ax1 = fig.add_subplot(111)

    x, y = zip(*sorted(zip(x, y)))

    ax1.plot(x, y, label="Logistic Regression")
    ax1.annotate(str(round(max(y), 4)), xy=(x[y.index(max(y))], y[y.
↪index(max(y))]),
                xytext=(x[y.index(max(y))], max(y) - 0.01),
                arrowprops=dict(facecolor='black', shrink=0.05))
    plt.xlabel('Hyperparameters')
    plt.ylabel('Validation Accuracy')
    plt.title('Result of Logistic Regression for different hyperparameter_
↪values')
    plt.grid(True)
    plt.savefig("../images/Final/logistic_regression_hyperparameter_tuning/lr.
↪png")
    plt.savefig("../images/Final/logistic_regression_hyperparameter_tuning/lr.
↪pdf")
    plt.show()

```



```
[ ]: with open("../model/Grid_search_modes/gs_lr_tuned.pkl", 'rb') as gs_lr_clf_file:
    gs_lr_clf = pickle.load(gs_lr_clf_file)

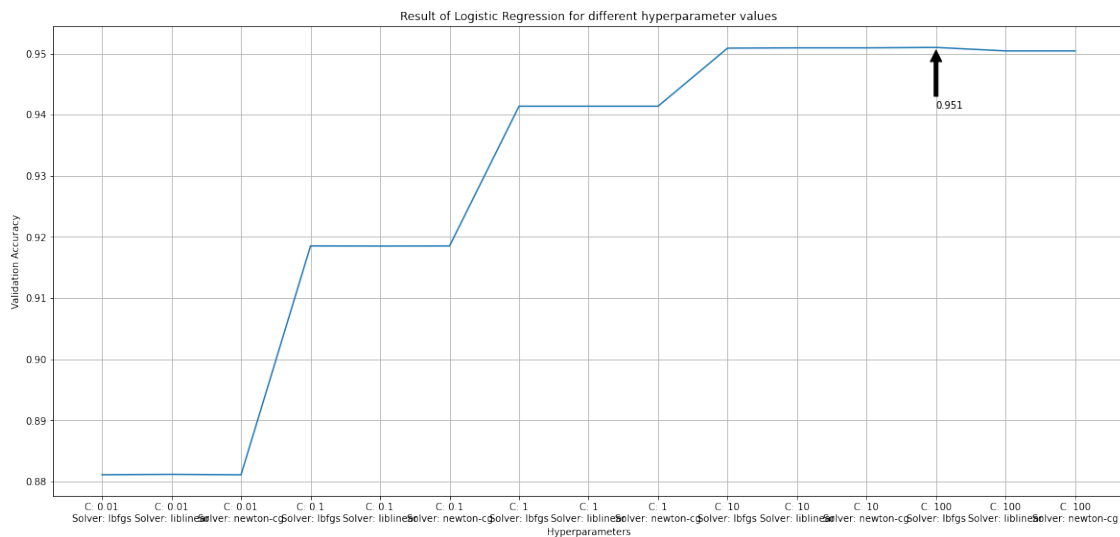
with open("../model/Grid_search_modes/gs_lr_features.pkl", 'rb') as gs_lr_features_file:
    gs_lr_features = pickle.load(gs_lr_features_file)

x = list()
params = gs_lr_clf.cv_results_['params']
for param in params:
    C = "C: " + str(param['clf__C'])
    solver = "Solver: " + str(param['clf__solver'])
    x.append(C + "\n" + solver)

lr_means = gs_lr_clf.cv_results_['mean_test_score']
y = lr_means

draw_hp_performance_lr(x, y)
lr_tuned_value = str("Best validation accuracy for logistic regression: %f using %s" % (max(y), x[y.index(max(y))]))

telegram_bot_sendtext(lr_tuned_value)
print("Best: %f using %s" % (max(y), x[y.index(max(y))]))
```



Best: 0.951012 using C: 100
 Solver: lbfgs

6.4 Tune hyperparameters of support vector machine

```
[ ]: estimators_svm = [  
    ('vect', CountVectorizer(ngram_range=(1,1))),  
    ('tfidf', TfidfTransformer(norm='l2')),  
    ('clf', SGDClassifier(tol=None, max_iter=100))  
]  
svm_clf = Pipeline(estimators_svm)  
  
param_grid = dict(  
    clf__alpha = [0.00001, 0.0001, 0.001, 0.01, 0.1],  
)  
# param_grid = {'clf__learning_rate': [0.1, 1, 10, 100],  
#               'clf__loss':['log']}  
  
kfold = KFold(n_splits=10, random_state=23)  
gs_svm_clf = GridSearchCV(svm_clf, param_grid=param_grid, n_jobs=-1, cv=kfold,  
    →scoring='accuracy')  
  
start = time.time()  
with warnings.catch_warnings():  
    warnings.filterwarnings("ignore")  
    gs_svm_clf = gs_svm_clf.fit(X_train, Y_train)  
end = time.time()  
print("Time taken:", end - start)  
  
svm_tuned = "Execution of hyperparameter tuning for svm is complete. Execution_  
    →time: "+str(end - start)+" sec"  
telegram_bot_sendtext(svm_tuned)  
  
with open("../model/Grid_search_modes/gs_svm_tuned.pkl", 'wb') as_  
    →gs_svm_clf_file:  
    pickle.dump(gs_svm_clf, gs_svm_clf_file, pickle.HIGHEST_PROTOCOL)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_split.py:296:  
FutureWarning: Setting a random_state has no effect since shuffle is False. This  
will raise an error in 0.24. You should leave random_state to its default  
(None), or set shuffle=True.
```

FutureWarning

Time taken: 613.0611274242401

```
[ ]: def draw_hp_performance_svm(x, y):  
    fig = plt.figure(figsize=(12,6))  
    ax1 = fig.add_subplot(111)  
  
    x, y = zip(*sorted(zip(x, y)))
```

```

ax1.plot(x, y, label="Support vector machine")
ax1.annotate(str(round(max(y), 4)), xy=(x[y.index(max(y))], y[y.
→index(max(y))]),
            xytext=(x[y.index(max(y))], max(y) - 0.01),
            arrowprops=dict(facecolor='black', shrink=0.05))
plt.xlabel('Hyperparameters')
plt.ylabel('Validation Accuracy')
plt.title('Result of support vector machine for different hyperparameter_
→values')
plt.grid(True)
plt.savefig("../images/Final/SVM_hyperparameter_tuning/svm.png")
plt.savefig("../images/Final/SVM_hyperparameter_tuning/svm.pdf")

plt.show()

```

```

[ ]: with open("../model/Grid_search_modes/gs_svm_tuned.pkl", 'rb') as
→gs_svm_clf_file:
    gs_svm_clf = pickle.load(gs_svm_clf_file)

with open("../model/Grid_search_modes/gs_svm_features.pkl", 'rb') as
→gs_svm_features_file:
    gs_svm_features = pickle.load(gs_svm_features_file)


x = list()
params = gs_svm_clf.cv_results_['params']
for param in params:
    alpha = "alpha: " + str(np.format_float_positional(param['clf__alpha'],
→trim='-'))
    x.append(alpha)

svm_means = gs_svm_clf.cv_results_['mean_test_score']
y = svm_means

draw_hp_performance_svm(x, y)
svm_tuned_value = str("Best accuracy for svm: %f using %s" % (max(y), x[y.
→index(max(y))]))

telegram_bot_sendtext(svm_tuned_value)
print("Best: %f using %s" % (max(y), x[y.index(max(y))]))

```



Best: 0.960713 using alpha: 0.00001

6.4.1 It is evident from the above result that support vector machine outperforms naive bayes and logistic regression for $\alpha = 0.00001$ (Learning rate), l2 norm and (1,1) n grams. So 6 different levels of toxicity is trained with svm classifier with above mentioned hyperparameter to achieve good result.

6.4.2 SVM classifier for 6 different classes

```
[7]: class SVM():
    def __init__(self, train_data, test_data, comment_type):
        if(comment_type not in train_data):
            raise Exception(comment_type + ' is not found in the training data')

        self.X_train = train_data['comment_text'].values
        self.Y_train = train_data[comment_type].values
        self.X_test = test_data['comment_text'].values
        self.Y_test = test_data[comment_type].values
        self.comment_type = comment_type

    def train(self):
        estimators_svm = [('vect', CountVectorizer(ngram_range=(1,1))),
                           ('tfidf', TfidfTransformer(norm='l2')),
                           ('clf', SGDClassifier(tol=None,
        ↪max_iter=100,alpha=0.00001))]

        self.svm_clf = Pipeline(estimators_svm)
```

```

print("start")
start = time.time()
with warnings.catch_warnings():
    warnings.filterwarnings("ignore")
    self.svm_clf = self.svm_clf.fit(self.X_train, self.Y_train)
end = time.time()
print("Time taken:", end - start)

svm_tuned_new = str(self.comment_type) + ": Training complete. Execution_
↪time: " + str(end - start) + " sec"
telegram_bot_sendtext(svm_tuned_new)

with open("../model/svm_trained_model/svm_" + str(self.comment_type) +
↪".pkl", 'wb') as svm_clf_file:
    pickle.dump(self.svm_clf, svm_clf_file, pickle.HIGHEST_PROTOCOL)

def load_model(self, file):
    self.svm_clf = pickle.load(open(file, 'rb'))

def performance_measure(self):
    self.prediction = self.svm_clf.predict(self.X_test)
    print(classification_report(self.prediction, self.Y_test))
    print("-"*20)
    print("confusion matrix")
    print(self.confusion_matrix())

def confusion_matrix(self):
    data = {'y_Actual': self.Y_test,
            'y_Predicted': self.prediction}

    df = pd.DataFrame(data, columns=['y_Actual', 'y_Predicted'])
    confusion_matrix = pd.crosstab(df['y_Actual'], df['y_Predicted'],
↪rownames=['Actual'], colnames=['Predicted'])
    sns.heatmap(confusion_matrix, annot=True, fmt='d')
    plt.savefig("../images/Final/Confusion_matrix/" + str(self.
↪comment_type) + ".pdf")
    plt.show()

```

6.4.3 Load data

```

[8]: train_data = pd.read_csv("../data/train.csv", index_col=False)

test_data_comments = pd.read_csv("../data/test.csv", index_col=False)
test_data_labels = pd.read_csv("../data/test_labels.csv", index_col=False)

```

```
test_data = pd.concat([test_data_comments, test_data_labels.drop(['id'],
↳axis=1)], axis=1, sort=False)
test_data = test_data[test_data['toxic'] != -1]
```

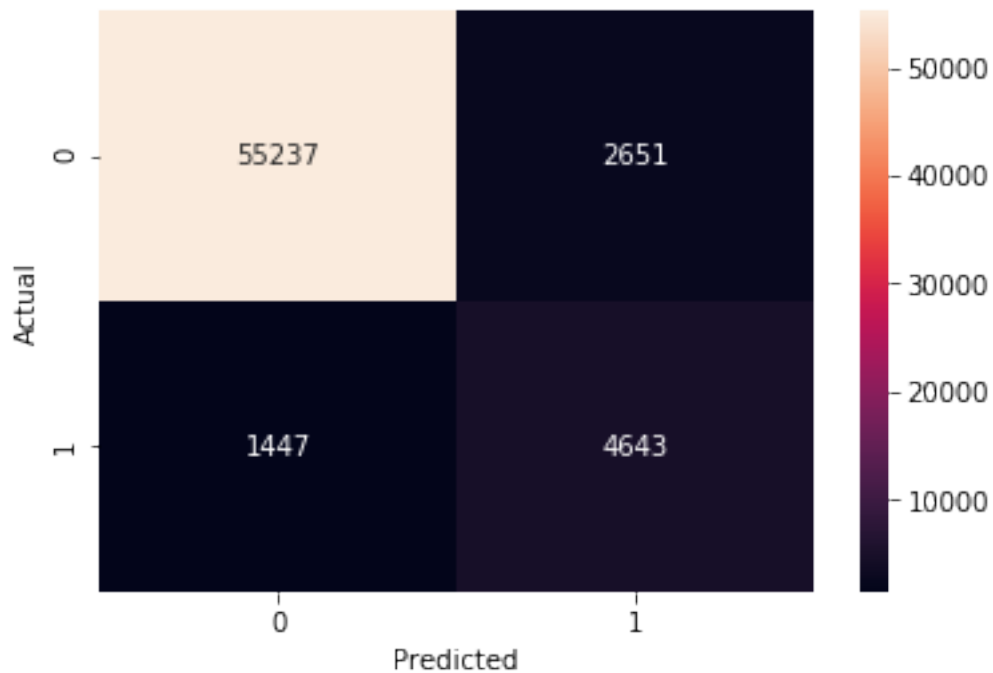
```
[9]: svm_toxic = SVM(train_data, test_data, 'toxic')
svm_toxic.train()
svm_toxic.performance_measure()
```

start

Time taken: 10.101313591003418

	precision	recall	f1-score	support
0	0.95	0.97	0.96	56684
1	0.76	0.64	0.69	7294
accuracy			0.94	63978
macro avg	0.86	0.81	0.83	63978
weighted avg	0.93	0.94	0.93	63978

confusion matrix



None


```
[10]: svm_severe_toxic = SVM(train_data, test_data, 'severe_toxic')
svm_severe_toxic.train()
svm_severe_toxic.performance_measure()
```

start

Time taken: 10.526662111282349

	precision	recall	f1-score	support
0	1.00	0.99	1.00	63911
1	0.07	0.39	0.12	67
accuracy			0.99	63978
macro avg	0.54	0.69	0.56	63978
weighted avg	1.00	0.99	1.00	63978

confusion matrix



None


```
[11]: svm_obscene = SVM(train_data, test_data, 'obscene')
svm_obscene.train()
svm_obscene.performance_measure()
```

start

Time taken: 12.23163652420044

	precision	recall	f1-score	support
0	0.98	0.98	0.98	60426
1	0.67	0.70	0.69	3552
accuracy			0.96	63978
macro avg	0.83	0.84	0.83	63978
weighted avg	0.97	0.96	0.96	63978

confusion matrix



None

```
[12]: svm_threat = SVM(train_data, test_data, 'threat')
      svm_threat.train()
      svm_threat.performance_measure()
```


start

Time taken: 12.167128324508667

	precision	recall	f1-score	support
0	1.00	1.00	1.00	63922
1	0.17	0.64	0.27	56

accuracy			1.00	63978
macro avg	0.59	0.82	0.63	63978
weighted avg	1.00	1.00	1.00	63978

confusion matrix



None

```
[13]: svm_insult = SVM(train_data, test_data, 'insult')
      svm_insult.train()
      svm_insult.performance_measure()
```


start

Time taken: 12.20309829711914

	precision	recall	f1-score	support
0	0.99	0.98	0.98	61214
1	0.56	0.70	0.63	2764

accuracy			0.96	63978
macro avg	0.78	0.84	0.80	63978
weighted avg	0.97	0.96	0.97	63978

confusion matrix



None


```
[14]: svm_identity_hate = SVM(train_data, test_data, 'identity_hate')
      svm_identity_hate.train()
      svm_identity_hate.performance_measure()
```

start

Time taken: 11.810347080230713

	precision	recall	f1-score	support
0	1.00	0.99	1.00	63780
1	0.21	0.76	0.33	198
accuracy			0.99	63978
macro avg	0.61	0.88	0.66	63978
weighted avg	1.00	0.99	0.99	63978

confusion matrix



None

6.5 Test and play

6.5.1 Load trained models

```
[34]: class Aggressive_classifier():

    def __init__(self, categories, models):
        if(len(categories) != len(models)):
            raise Exception("The number of categories is not the same as the_
↪number of models")
        self.categories = categories
        self.models = models

    def predict(self, sentence):
        for category, model in zip(self.categories, self.models):
            prediction = model.predict( [sentence] )

            if(prediction == 0):
                result = "no"
            else:
                result = "yes"

            print(category + " : " + result)
```

```
[35]: with open("../model/svm_trained_model/svm_toxic.pkl", 'rb') as f:
    gs_nb_features_file:
        svm_toxic = pickle.load(gs_nb_features_file)

with open("../model/svm_trained_model/svm_severe_toxic.pkl", 'rb') as f:
    gs_nb_features_file:
        svm_severe_toxic = pickle.load(gs_nb_features_file)

with open("../model/svm_trained_model/svm_obscene.pkl", 'rb') as f:
    gs_nb_features_file:
        svm_obscene = pickle.load(gs_nb_features_file)

with open("../model/svm_trained_model/svm_threat.pkl", 'rb') as f:
    gs_nb_features_file:
        svm_threat = pickle.load(gs_nb_features_file)

with open("../model/svm_trained_model/svm_insult.pkl", 'rb') as f:
    gs_nb_features_file:
        svm_insult = pickle.load(gs_nb_features_file)

with open("../model/svm_trained_model/svm_identity_hate.pkl", 'rb') as f:
    gs_nb_features_file:
        svm_identity_hate = pickle.load(gs_nb_features_file)

my_classifier = Aggressive_classifier(["toxic", "severe_toxic", "obscene",
    "threat", "insult", "identity_hate"],
    [svm_toxic, svm_severe_toxic, svm_obscene,
    svm_threat, svm_insult, svm_identity_hate])
```

6.5.2 Test models with some tweets

Test with sentences found in internet

```
[70]: sentence = "Thats not what you mum said when I was fucking her from behind last_
    night like the bitch she is."
print(sentence)
my_classifier.predict(sentence)
print()

sentence = "Oooh look at a little cry baby, aren't you just the sorriest lil'_
    bitch that ever breathed."
print(sentence)
my_classifier.predict(sentence)
print()

sentence = "Why don't you go eat a nigger you fat bastard."
print(sentence)
```

```

my_classifier.predict(sentence)
print()

sentence = "I have scraped shit off my shoe that looks prettier than you."
print(sentence)
my_classifier.predict(sentence)
print()

sentence = "Why don't you just fuck off and die, dick-bag."
print(sentence)
my_classifier.predict(sentence)
print()

sentence = "You're face looks like you were drop kicked as a kid into a food_
↳blender."
print(sentence)
my_classifier.predict(sentence)
print()

sentence = "Yo mumma so fat that is took 3 moutnainteering teams to find your_
↳ugly mug when you were born."
print(sentence)
my_classifier.predict(sentence)
print()

```

Thats not what you mum said when I was fucking her from behind last night like the bitch she is.

```

toxic : yes
severe_toxic : no
obscene : yes
threat : no
insult : yes
identity_hate : no

```

Oooh look at a little cry baby, aren't you just the sorriest lil' bitch that ever breathed.

```

toxic : yes
severe_toxic : no
obscene : yes
threat : no
insult : yes
identity_hate : no

```

Why don't you go eat a nigger you fat bastard.

```

toxic : yes
severe_toxic : no
obscene : yes

```

```
threat : no
insult : yes
identity_hate : yes
```

I have scraped shit off my shoe that looks prettier than you.

```
toxic : yes
severe_toxic : no
obscene : yes
threat : no
insult : no
identity_hate : no
```

Why don't you just fuck off and die, dick-bag.

```
toxic : yes
severe_toxic : no
obscene : yes
threat : no
insult : yes
identity_hate : no
```

You're face looks like you were drop kicked as a kid into a food blender.

```
toxic : no
severe_toxic : no
obscene : no
threat : no
insult : no
identity_hate : no
```

Yo mumma so fat that it took 3 mountainteering teams to find your ugly mug when you were born.

```
toxic : yes
severe_toxic : no
obscene : no
threat : no
insult : yes
identity_hate : no
```

Implicit threats are not detected

```
[64]: implicit_threat = "Any difficulty and we will assume control but, when the_
    ↳looting starts, the shooting starts. Thank you!"

print(implicit_threat)
my_classifier.predict(implicit_threat)
```

Any difficulty and we will assume control but, when the looting starts, the shooting starts. Thank you!

```
toxic : no
severe_toxic : no
obscene : no
threat : no
insult : no
identity_hate : no
```

Verb conjugation not detected

```
[65]: verb_infinity = "I kill people"
      print(verb_infinity)
      my_classifier.predict(verb_infinity)

      print()

      verb_gerundio = "I am killing people"
      print(verb_gerundio)
      my_classifier.predict(verb_gerundio)
```

```
I kill people
toxic : yes
severe_toxic : no
obscene : no
threat : yes
insult : no
identity_hate : no
```

```
I am killing people
toxic : no
severe_toxic : no
obscene : no
threat : no
insult : no
identity_hate : no
```

Reference - <https://www.kaggle.com/c/jigsaw-multilingual-toxic-comment-classification/data>

```
[ ]:
```